


Reinforcement Learning



(AI)

Table of contents

- 01** —→ Prerequisites and revisiting Neural Networks
- 02** —→ A basic investment strategy
- 03** —→ Reinforcement Learning: Core concepts
- 04** —→ Reinforcement Learning: Q-Learning

01 →

Prerequisites and revisiting Neural Networks

(AI)

Requirements for this seminar

Knowledge	Origin	Relevance
Deep Neural Networks	Deep Learning	High
Gradient Descent	Deep Learning	Medium
Cost Function	Data Analysis	Medium
Keras/Tensorflow	Deep Learning	High
R Programming	Data Analysis	High
Math Skills	School/Uni	Low
Knowledge Transfer	Uni	High



Organization

In week one you will receive a **brief introduction** to the current topic. Afterwards, you must make a **binding registration** for the seminar until Sunday, the 08.06.2025 23:59, via mail to steinert@europa-uni.de. This registration always leads to a grade.

In weeks 3, 4 and 5 there will be three preparation tasks you need to solve leading to **mandatory presentations** on the group level including graded questions by me. Weeks six to eight will be independent group work with individual **meetings with me on demand**. The submission deadline for the fourth and **final task** is Sunday, 27.07.2025, at 23:59.




Your seminar project

- Create a standard NN which aims to learn the optimal investment strategy
- Build the environment and replay buffer for reinforcement learning
- Build and train a Q-learning NN architecture
- Design and train your own Q-Learning NN architecture



Simple Neural Network (NN)

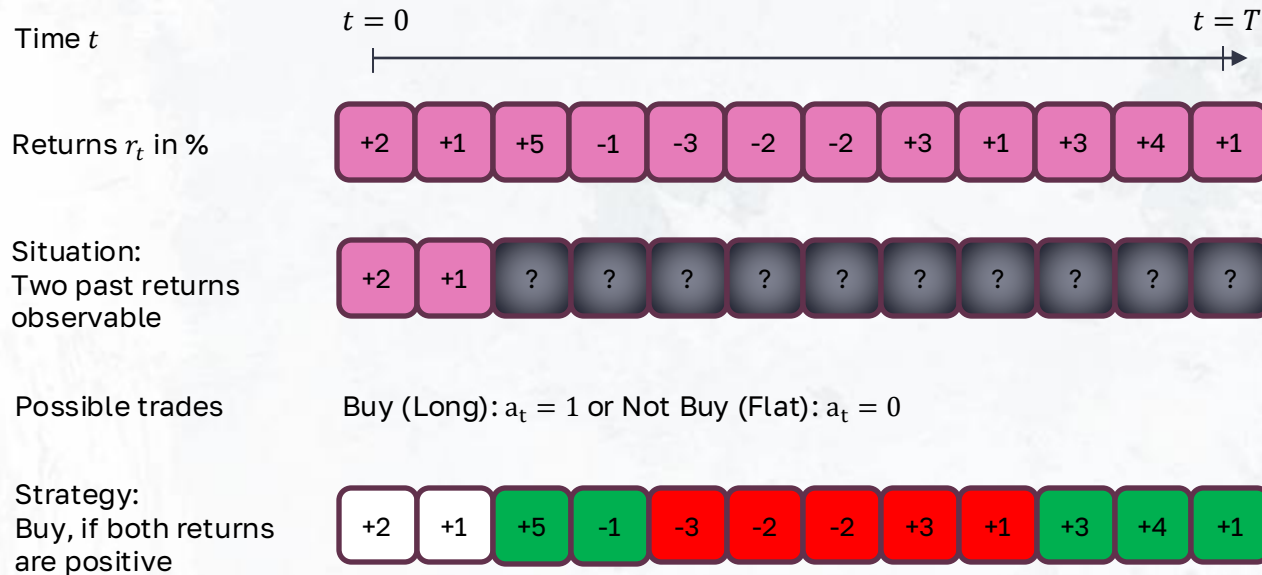


02 →

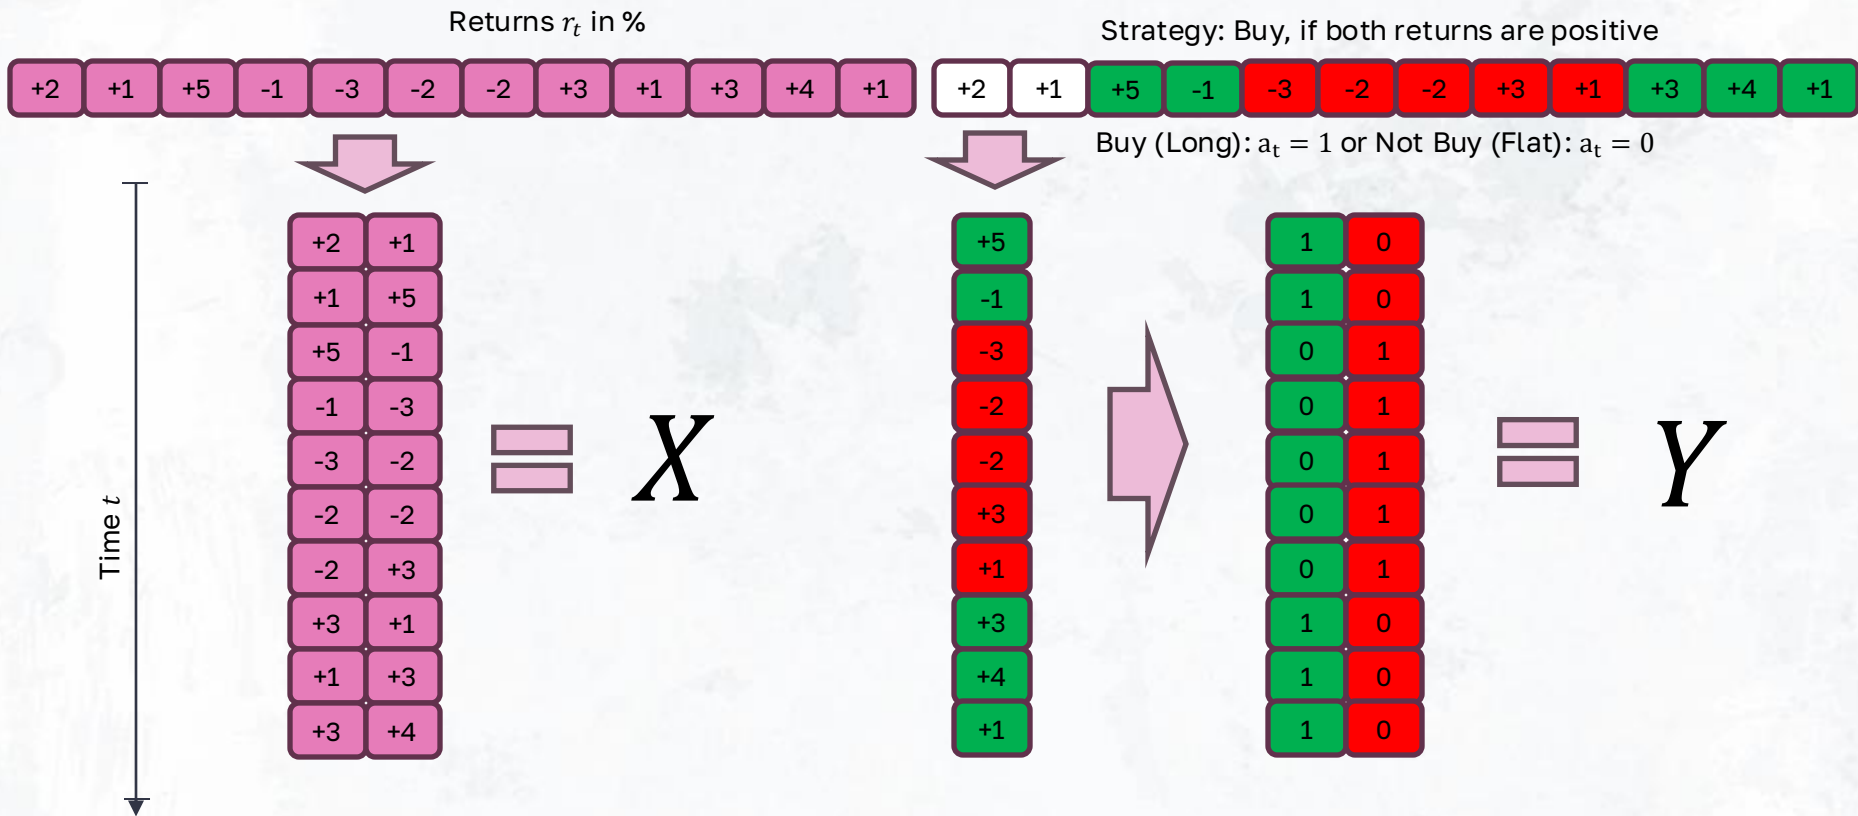
A basic investment strategy

(AI)

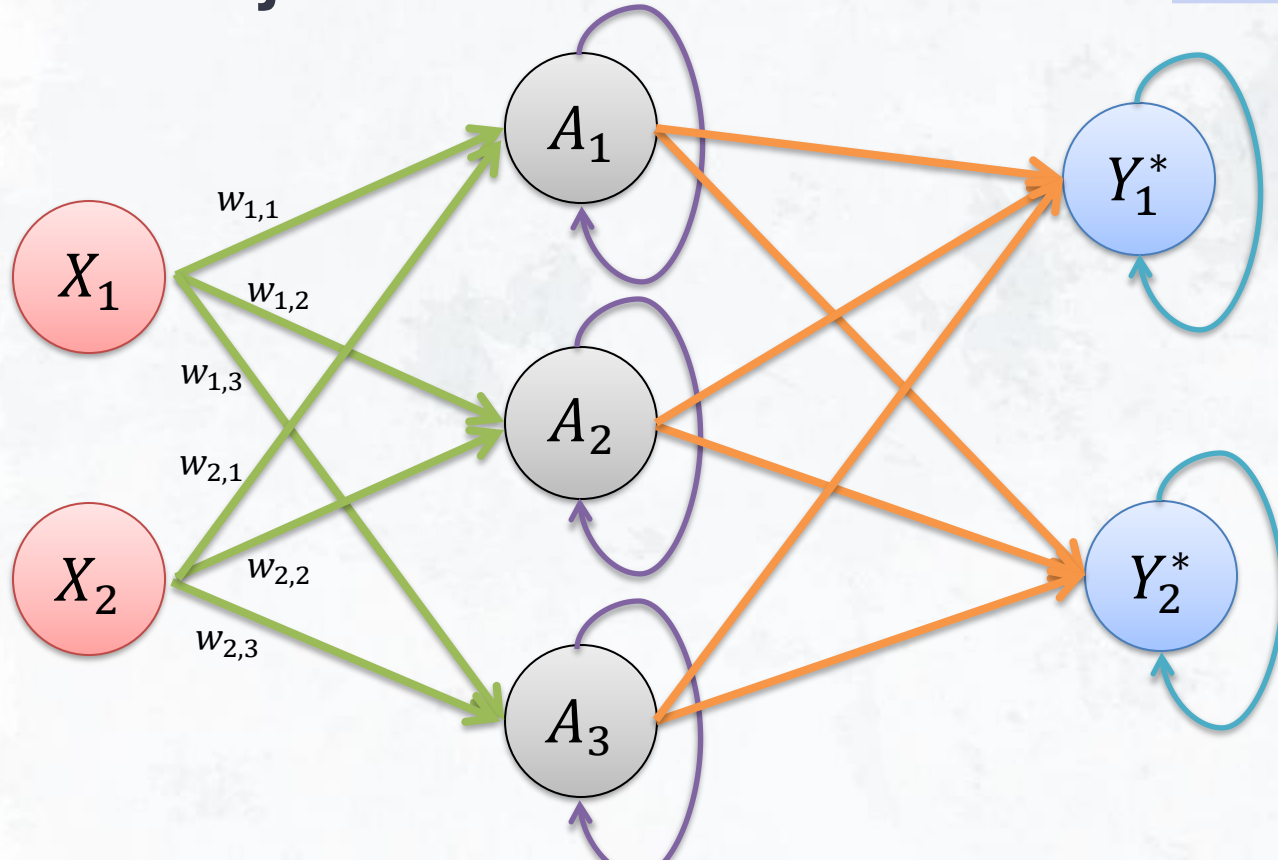
To buy or not to buy – that's the question!




From trades to supervised learning



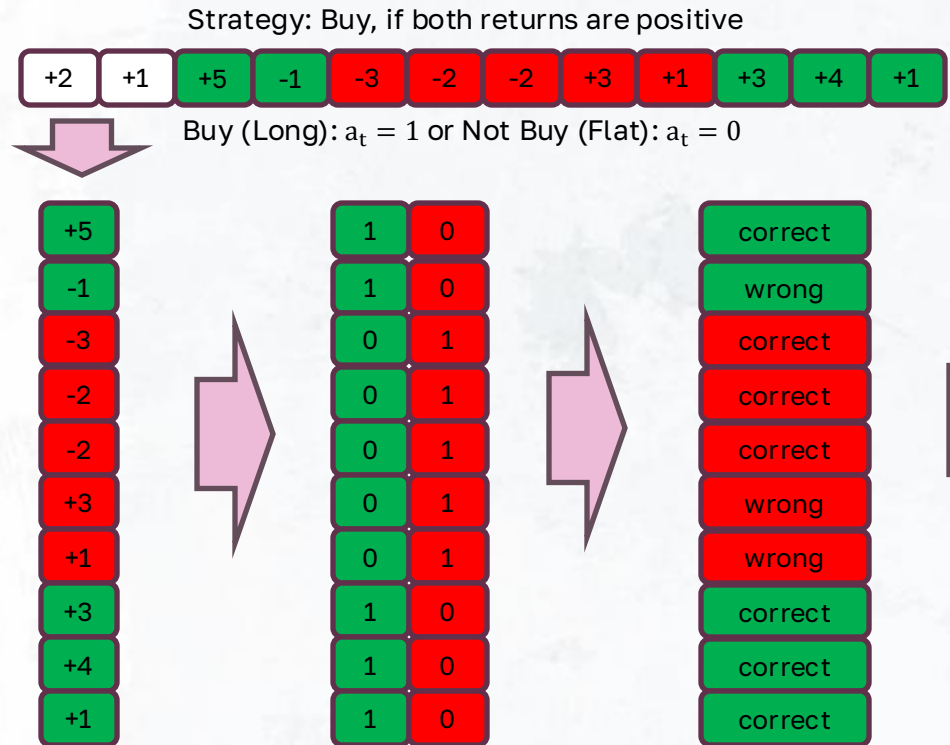
Idea: Now just model it with NN!



Let us focus, what does this neural network **actually** learn?

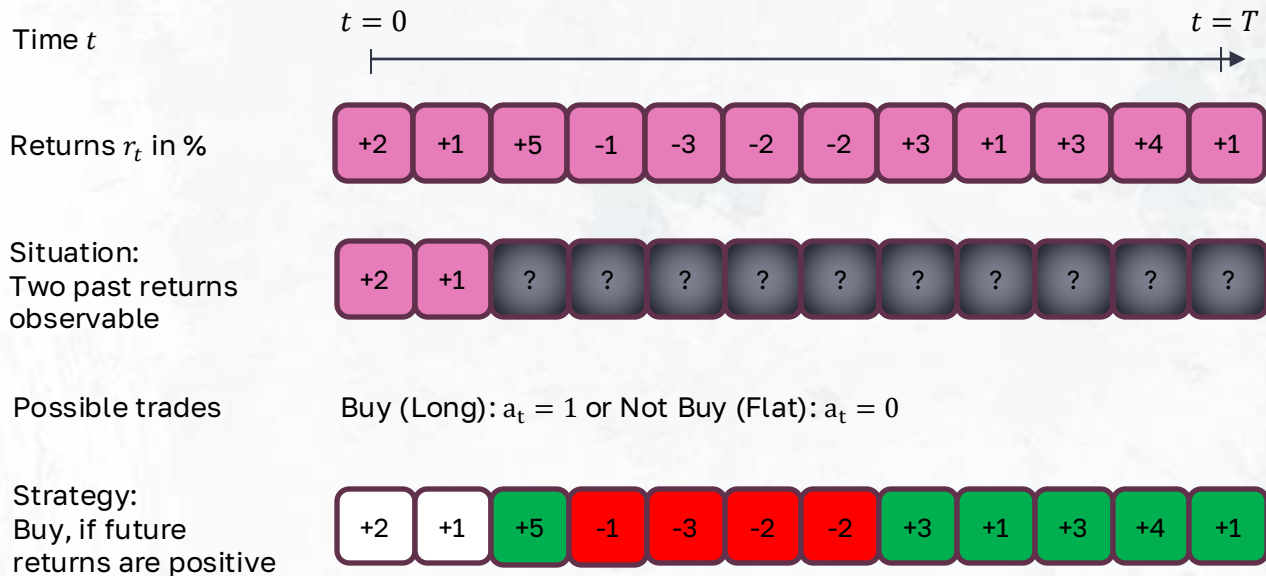


How good was this strategy?

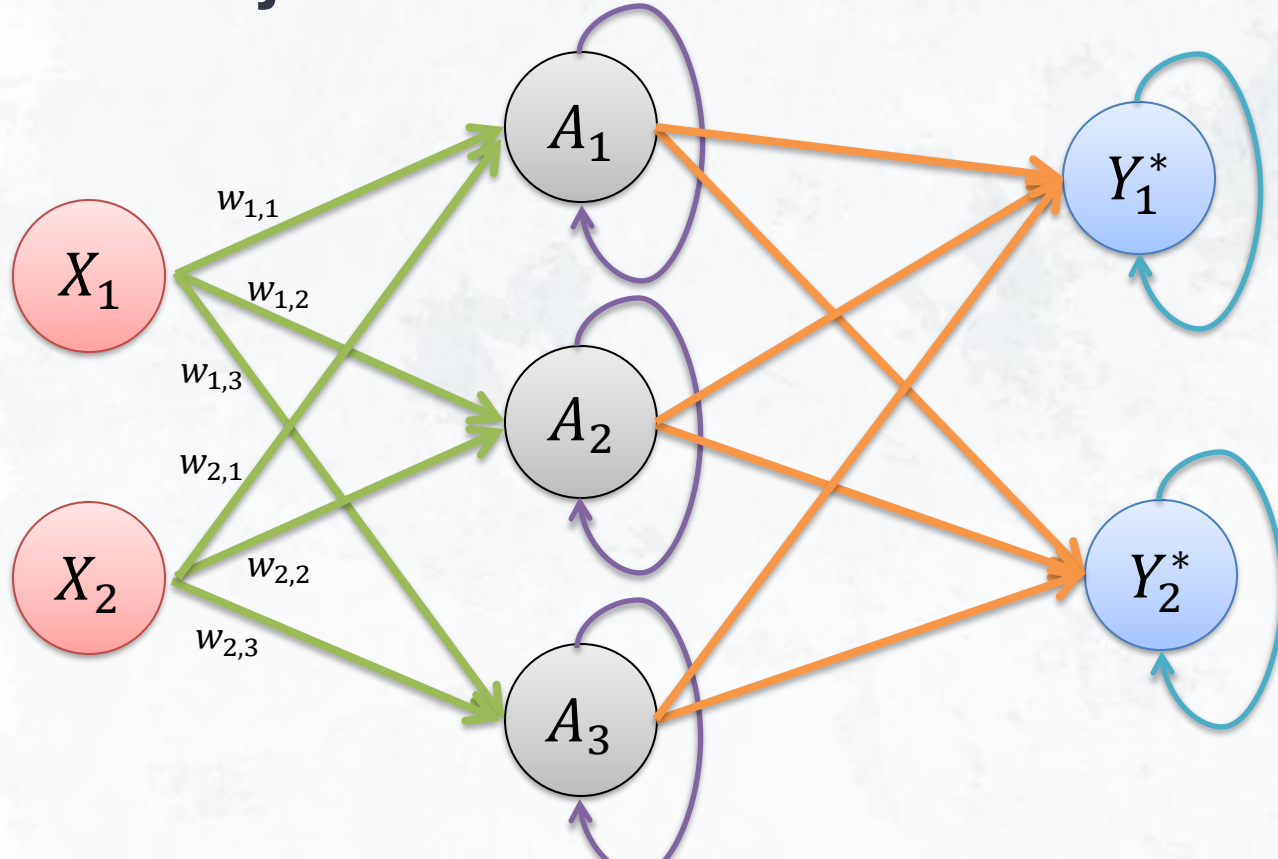


The strategy was imperfect, so we need to learn from other (better) traders!

Idea: Just buy, when it goes up!




Idea: Now just model it with NN!



How good was this strategy?

Let us focus, can a neural network actually learn from this?



Strategy: Buy, if future returns are positive

+2	+1	+5	-1	-3	-2	-2	+3	+1	+3	+4	+1
----	----	----	----	----	----	----	----	----	----	----	----



Buy (Long): $a_t = 1$ or Not Buy (Flat): $a_t = 0$

+5
-1
-3
-2
-2
+3
+1
+3
+4
+1



1	0
0	1
0	1
0	1
0	1
1	0
1	0
1	0
1	0
1	0



correct
correct
correct
correct
correct
correct
correct
correct
correct
correct



The strategy was perfect, so we are done, right?

Current Issues

We only modeled one trader

- In order to model the whole universe of strategies, we would need to observe all possible trades. This would amount to 2^T strategies!

We tried to model the perfect scenario

- We used the future return to retrieve the best outcome, but this is not observable and thus not learnable in real life.

We looked at results in the next time period

- Buying a stock gives the next days' return, but also comes with consequences later on, e.g. the need to sell, which involves transaction costs.

03 →

Reinforcement Learning: Core concepts

(AI)

Key idea of Reinforcement Learning

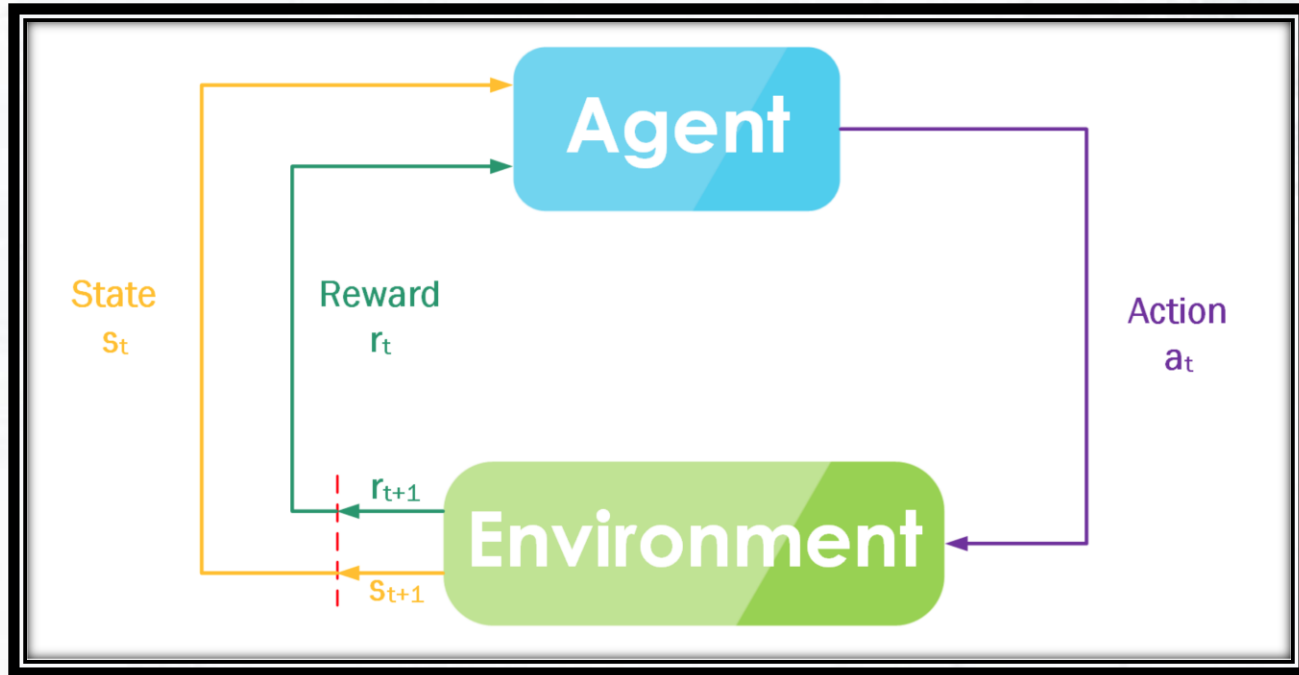
The main idea of Reinforcement Learning (RL) is to teach the AI to **interact** with the surrounding **environment**. By doing so, it **receives rewards**, which teaches it to improve.

In particular, it balances between trying out (**exploration**) and earning from what it has learned (**exploitation**). Just like humans, the AI needs to learn from the mistakes it makes as much as from the successes.

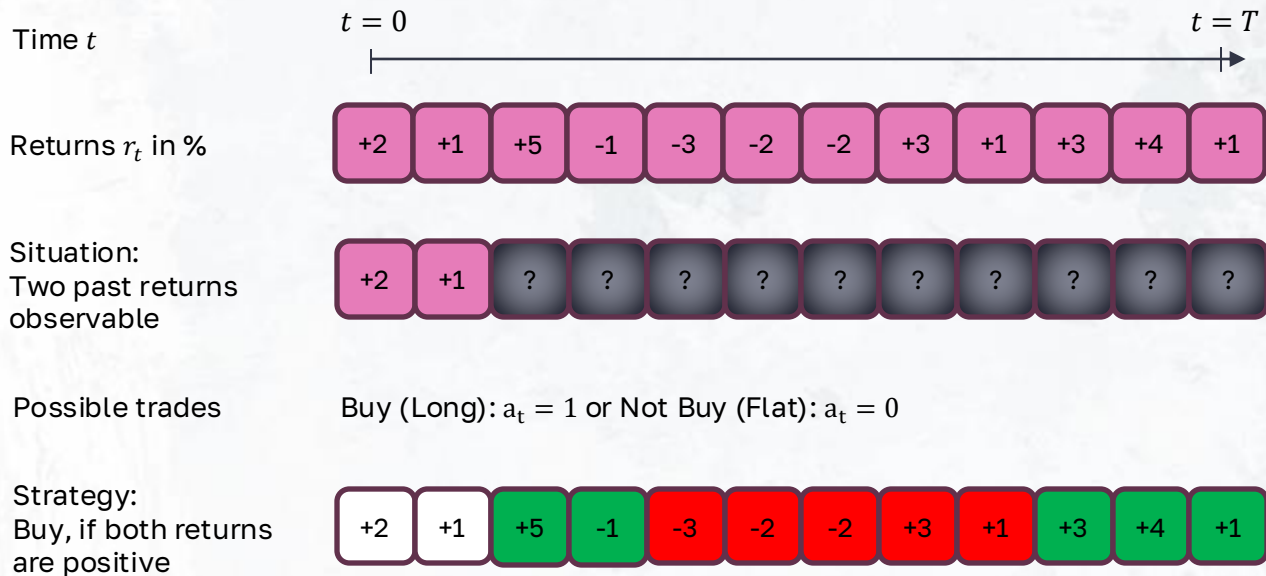
RL is applied in various fields, from robotics over trading to self-driving and gaming you will find a lot of implementations.



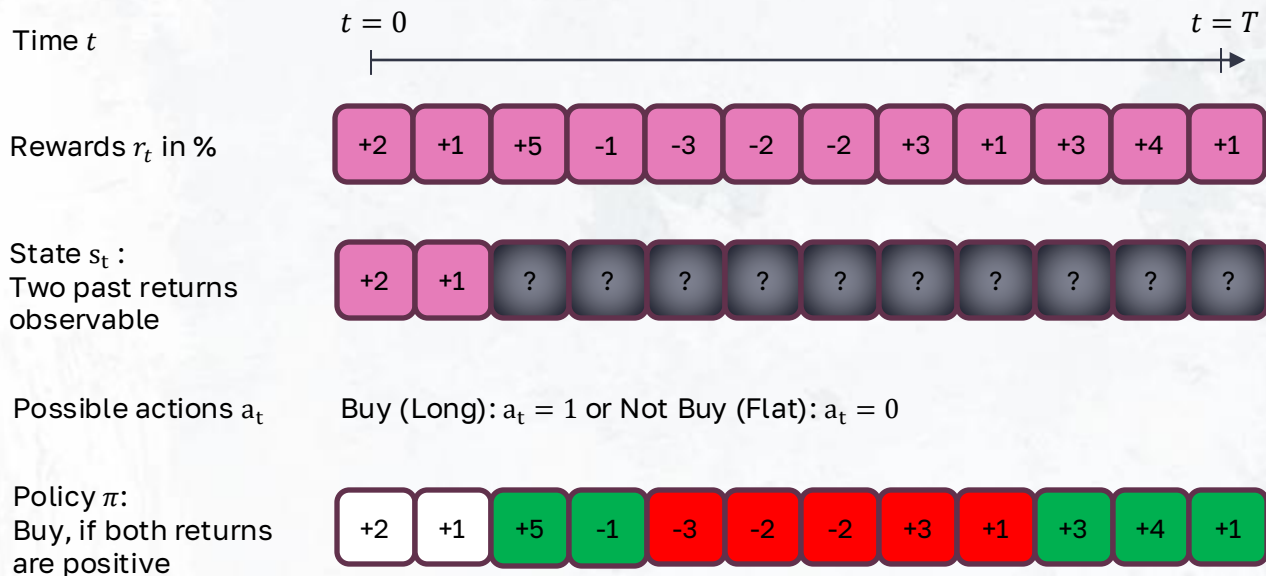
The core loop of RL



Remember this slide? Let's update!



Remember this slide? Let's update!



Overview of core elements in RL

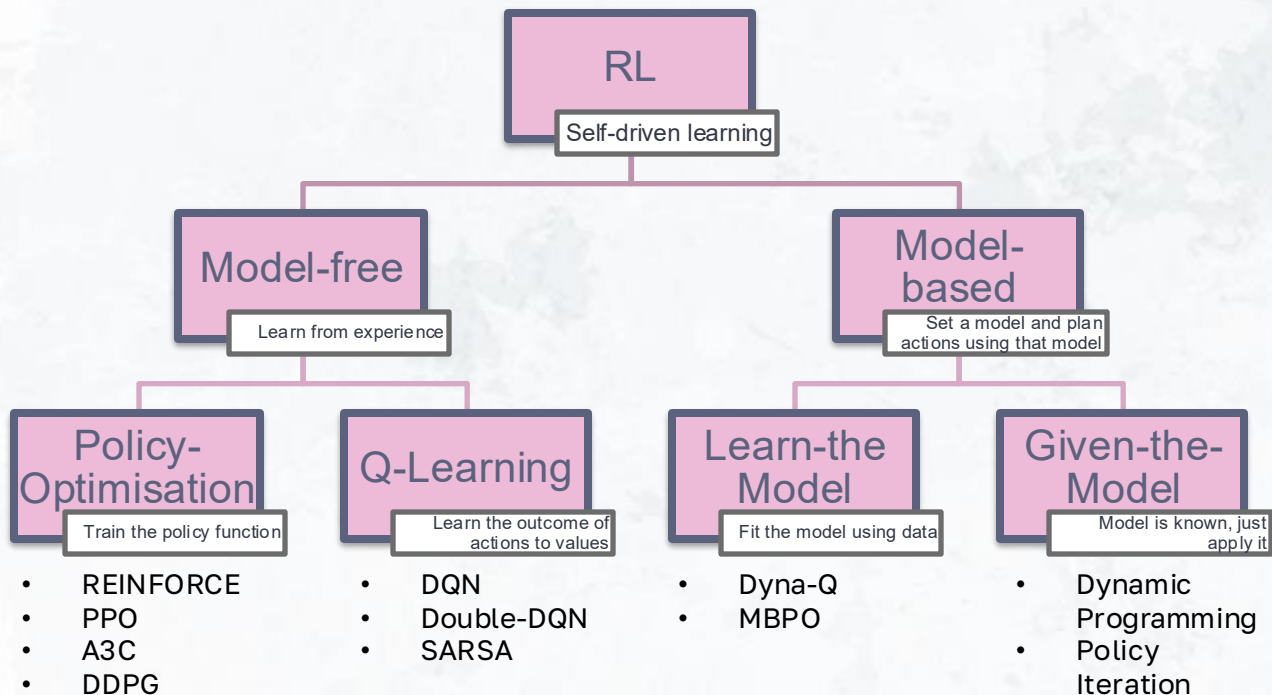
Term	Meaning general	Meaning in our example
Agent	Decision-making entity	Investor
Environment	Setting in which an agent operates	Daily stock market with prices and returns
State s	Representation of the current situation in the environment	Last $k = 2$ returns
Action a	A decision an agent can make	Long (1) or Flat (0)
Reward r	A feedback signal after taking an action	Return on investment
Policy π	Mapping from states to actions	Investment strategy, e.g. long, if both returns positive

04 →

Reinforcement Learning: Q-Learning

(AI)

Overview of RL frameworks



Q-Learning: DQN

$$Q^*(s, a) = E \left[\sum_{k=0}^{\infty} \gamma^k r_{t+k+1} \mid s_t = s, a_t = a \right]$$

- Goal: Find parameters θ , such that $Q_{\theta}(s, a) \approx Q^*(s, a)$; γ is a discount factor
- $Q^*(s, a)$ is the optimal action-value function, as it perfectly maps actions to states
- We achieve that, by optimizing (using simple MSE):

$$L(\theta) = E_{(s,a,r,s') \sim D} [(y - Q_{\theta}(s, a))^2]$$

- Where we create the target y as:

$$y = r + \gamma \max_{a'} Q_{\bar{\theta}}(s', a')$$

} Bellman Target

- The Q-function is essentially our NN, but we use it to create targets and estimate!
- To achieve this, we sporadically update, so that $Q_{\theta}(s, a) = Q_{\bar{\theta}}(s, a)$

Q-Learning: The Magic

Is the dog catching it's own tail?

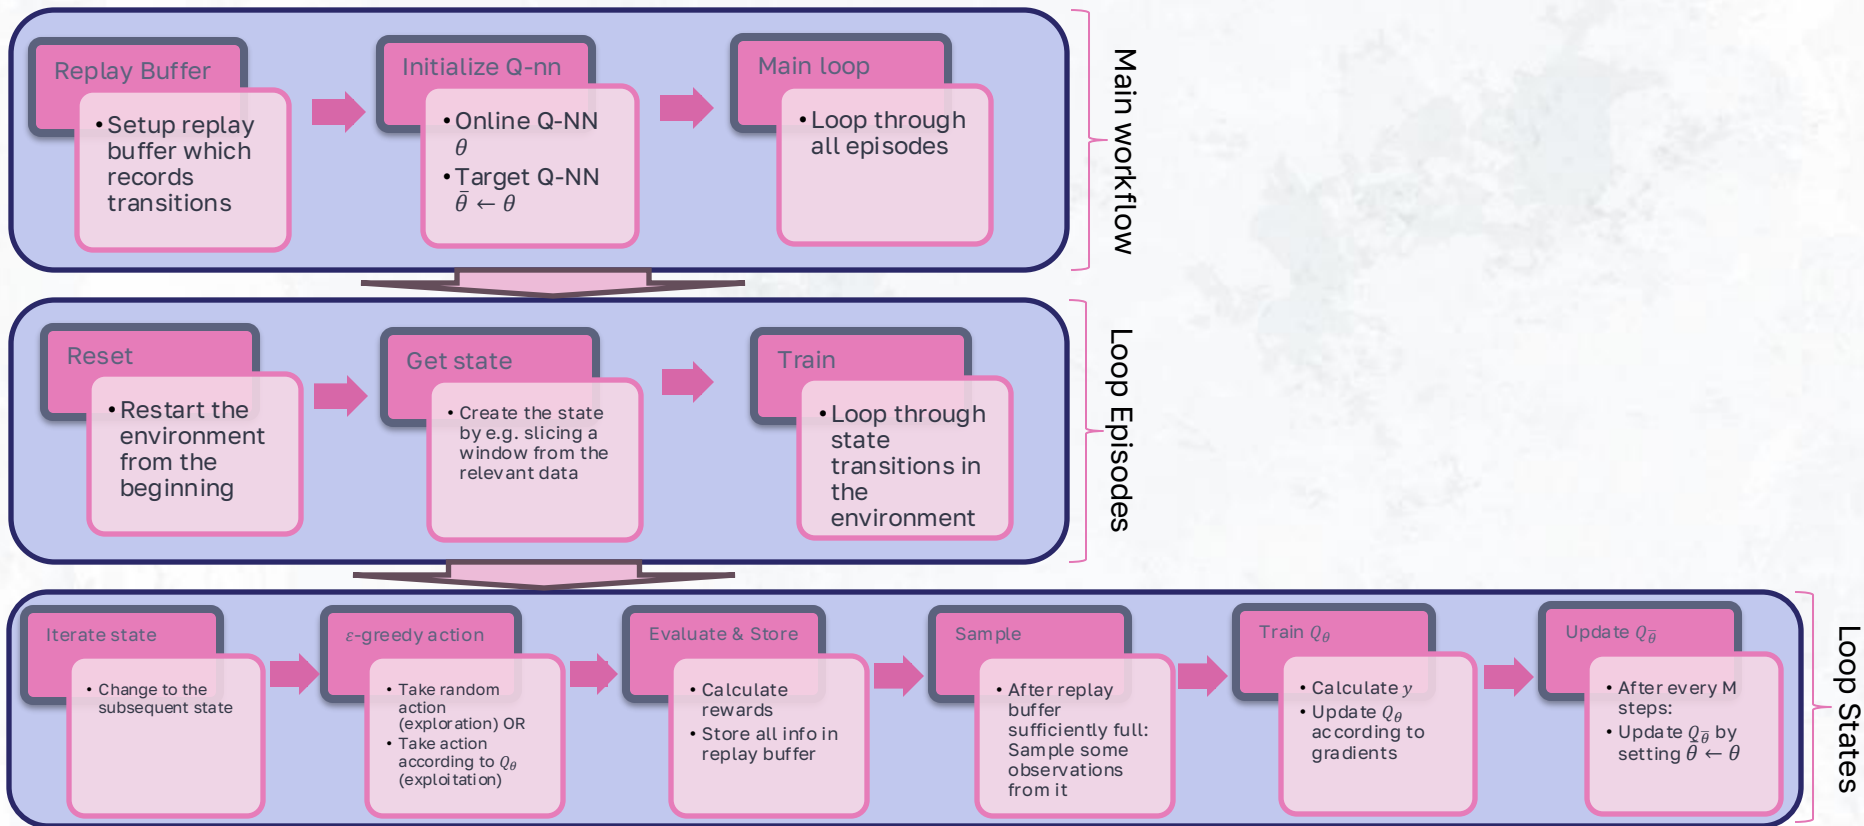
- OR -

How can the trainable network learn from targets generated by itself?

In simple terms:

- Training and target network start identical, but we freeze the target network temporarily
- Targets created contain the **reward r** , i.e. the networks are not purely self referential
- When taking actions, we either use the action suggested by the training network (**exploitation**) or a random one (**exploration**)
- Taking random actions helps us to learn what happens alternatively
- We **record** all of those actions and then **learn from the consequences** (gradient descent)
- Finally, we copy the weights of the training network into the target network and keep going

Simplified workflow of DQN



Thanks! →

Any questions?

Dr. Rick Steinert

steinert@europa-uni.de

CREDITS: This presentation template was created by **Slidesgo** and includes icons by **Flaticon**, infographics & images by **Freepik** and content by **Eliana Delacour**

